

# PROJECT-BASED AI SKILL-UP ROADMAP

## Learn by Building: "IndustrialMind" - End-to-End ML Platform

Diego Carriel Lopez | 12-Month Project Journey

### ⌚ THE PHILOSOPHY

"The best way to learn is to build something real."

Instead of: Learn PyTorch → Learn MLOps → Learn K8s → Build Project

We do: Build Project → Learn PyTorch WHILE building → Learn MLOps WHILE deploying → etc.

Every week you ship something. Every feature teaches a skill.

### ⚙️ THE PROJECT: IndustrialMind

#### What We're Building

An **open-source Industrial AI Platform** that: 1. Ingests real-time sensor data from manufacturing equipment 2. Detects anomalies and predicts failures using ML 3. Provides a Knowledge Graph of equipment relationships 4. Offers an LLM-powered assistant for technicians 5. Deploys with full MLOps pipeline

#### Why This Project?

Your Experience	Project Component	Skills Gained
Nestlé InfluxDB work	Data ingestion layer	Real-world relevance
Neo4j certification	Knowledge Graph module	Showcase expertise
LangChain/RAG experience	LLM Assistant	Deepen LLM skills
Streamlit apps	Dashboard	Production UI
<b>NEW</b>	PyTorch models	Fill critical gap
<b>NEW</b>	MLOps pipeline	Fill critical gap
<b>NEW</b>	Kubernetes deployment	Fill critical gap

#### The Stack (Matches \$200K+ Job Requirements)



# 12-MONTH PROJECT ROADMAP

## Overview

Month	Project Phase	Key Deliverable	Skills Acquired
1	Project Setup & Data Simulator	Working data pipeline	Kafka, Docker basics
2	First PyTorch Model	Anomaly detector	PyTorch fundamentals
3	MLflow Integration	Tracked experiments	MLOps basics
4	Time Series Forecasting	Predictive model	Advanced PyTorch
5	Knowledge Graph	Equipment graph	Neo4J advanced
6	RAG System	Document Q&A	Vector DBs, LangChain
7	LLM Fine-tuning	Domain-adapted LLM	LoRA, PEFT
8	Kubernetes Deployment	K8s cluster	K8s, Helm
9	Cloud Migration	AWS/Azure deployment	Cloud ML platforms
10	CI/CD Pipeline	Automated deployment	GitHub Actions, MLOps
11	Monitoring & Observability	Production monitoring	Prometheus, Grafana
12	Polish & Documentation	Portfolio-ready	Technical writing

## MONTH 1: Foundation & Data Pipeline

**Goal:** Build a working data ingestion system that simulates industrial sensors

### Week 1: Project Bootstrap

#### Tasks:

```
# Day 1-2: Repository setup
mkdir industrialmind && cd industrialmind
git init
# Create structure:
# industrialmind/
#   ├── docker-compose.yml
#   ├── README.md
#   ├── data-simulator/
#   ├── data-ingestion/
#   ├── ml-models/
#   ├── knowledge-graph/
#   ├── llm-assistant/
#   ├── api/
#   ├── frontend/
#   └── mlops/
#     └── docs/
```

**Deliverable:** GitHub repo with proper structure, README, and .gitignore

**Learn While Building:** - Resource: [Docker in 1 Hour](#) (FREE) - Resource: [Git Best Practices](#)

### Week 2: Industrial Data Simulator

**Tasks:** - Build Python simulator that generates realistic sensor data - Simulate: temperature, vibration, pressure, power consumption - Add realistic patterns: normal operation, degradation, failure modes

#### Code to Write:

```
# data-simulator/simulator.py
import numpy as np
from dataclasses import dataclass
from enum import Enum

class MachineState(Enum):
    NORMAL = "normal"
    DEGRADING = "degrading"
    FAILING = "failing"

@dataclass
class SensorReading:
    timestamp: datetime
    machine_id: str
    temperature: float
    vibration: float
    pressure: float
    power: float
    state: MachineState # for training labels

class IndustrialSimulator:
```

```

def generate_reading(self, machine_id: str) -> SensorReading:
    # Implement realistic sensor patterns
    pass

```

**Deliverable:** Working simulator producing 1000+ readings/minute

**Learn While Building:** - Resource: [Time Series Simulation](#)

### Week 3: Kafka + InfluxDB Pipeline

**Tasks:** - Set up Kafka for streaming (Docker) - Set up InfluxDB for storage (Docker) - Connect simulator → Kafka → InfluxDB

**docker-compose.yml:**

```

version: '3.8'
services:
  zookeeper:
    image: confluentinc/cp-zookeeper:latest
    environment:
      ZOOKEEPER_CLIENT_PORT: 2181

  kafka:
    image: confluentinc/cp-kafka:latest
    depends_on:
      - zookeeper
    ports:
      - "9092:9092"
    environment:
      KAFKA_BROKER_ID: 1
      KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://localhost:9092

  influxdb:
    image: influxdb:2.7
    ports:
      - "8086:8086"
    volumes:
      - influxdb-data:/var/lib/influxdb2

  simulator:
    build: ./data-simulator
    depends_on:
      - kafka

  ingestion:
    build: ./data-ingestion
    depends_on:
      - kafka
      - influxdb

volumes:
  influxdb-data:

```

**Deliverable:** docker-compose up starts entire pipeline

**Learn While Building:** - Resource: [Kafka Basics in 30 min](#) (FREE) - Resource: [InfluxDB Python Client](#)

### Week 4: Basic Visualization

**Tasks:** - Create simple Streamlit dashboard - Show real-time sensor readings - Plot time series with Plotly

**Deliverable:** Live dashboard at <http://localhost:8501>

**Learn While Building:** - You already know Streamlit! Just apply it.

### Month 1 Checkpoint ✓

- GitHub repo with proper structure
- Docker Compose running all services
- Simulator generating realistic data
- Kafka streaming working
- InfluxDB storing data
- Basic Streamlit dashboard
- README with setup instructions

**Skills Acquired:** - ✓ Docker & Docker Compose - ✓ Kafka fundamentals - ✓ Streaming data architecture - ✓ Project organization

## MONTH 2: First PyTorch Model

### Goal: Build an anomaly detection model in PyTorch (NOT TensorFlow!)

#### Week 5: PyTorch Fundamentals

**Tasks:** - Convert your existing TensorFlow knowledge to PyTorch - Build a simple autoencoder for anomaly detection

**Learn While Building:**

```

# ml-models/anomaly_detector/model.py
import torch
import torch.nn as nn

class SensorAutoencoder(nn.Module):
    def __init__(self, input_dim: int = 4, latent_dim: int = 2):
        super().__init__()

        # Encoder
        self.encoder = nn.Sequential(
            nn.Linear(input_dim, 16),
            nn.ReLU(),
            nn.Linear(16, 8),
            nn.ReLU(),
            nn.Linear(8, latent_dim)
        )

        # Decoder
        self.decoder = nn.Sequential(
            nn.Linear(latent_dim, 8),
            nn.ReLU(),
            nn.Linear(8, 16),
            nn.ReLU(),
            nn.Linear(16, input_dim)
        )

    def forward(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded

    def get_reconstruction_error(self, x):
        reconstructed = self.forward(x)
        return torch.mean((x - reconstructed) ** 2, dim=1)

```

**Resource:** [PyTorch in 60 Minutes](#) (FREE, Official)

## Week 6: Training Pipeline

**Tasks:** - Create PyTorch Dataset and DataLoader - Implement training loop with validation - Add early stopping

**Code Structure:**

```

# ml-models/anomaly_detector/train.py
from torch.utils.data import Dataset, DataLoader

class SensorDataset(Dataset):
    def __init__(self, influxdb_client, time_range):
        self.data = self.load_from_influx(influxdb_client, time_range)

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        return torch.tensor(self.data[idx], dtype=torch.float32)

def train_model(model, train_loader, val_loader, epochs=100):
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
    criterion = nn.MSELoss()

    for epoch in range(epochs):
        model.train()
        for batch in train_loader:
            optimizer.zero_grad()
            output = model(batch)
            loss = criterion(output, batch)
            loss.backward()
            optimizer.step()

        # Validation
        model.eval()
        val_loss = evaluate(model, val_loader, criterion)
        print(f"Epoch {epoch}: Val Loss = {val_loss:.4f}")

```

**Deliverable:** Trained model with validation metrics

## Week 7: Model Evaluation & Thresholding

**Tasks:** - Implement anomaly scoring - Find optimal threshold using validation data - Create evaluation metrics (precision, recall, F1)

**Code:**

```

# ml-models/anomaly_detector/evaluate.py
import numpy as np
from sklearn.metrics import precision_recall_curve, f1_score

def find_optimal_threshold(model, val_data, val_labels):
    """Find threshold that maximizes F1 score"""

```

```

model.eval()
with torch.no_grad():
    errors = model.get_reconstruction_error(val_data).numpy()

precisions, recalls, thresholds = precision_recall_curve(val_labels, errors)
f1_scores = 2 * (precisions * recalls) / (precisions + recalls + 1e-8)

optimal_idx = np.argmax(f1_scores)
return thresholds[optimal_idx], f1_scores[optimal_idx]

```

## Week 8: Real-time Inference Service

**Tasks:** - Create FastAPI endpoint for predictions - Dockerize the model service - Connect to Kafka for real-time scoring

**Code:**

```

# api/anomaly_service.py
from fastapi import FastAPI
import torch

app = FastAPI()

# Load model at startup
model = SensorAutoencoder()
model.load_state_dict(torch.load("models/autoencoder.pt"))
model.eval()
threshold = 0.05 # From evaluation

@app.post("/predict")
async def predict_anomaly(reading: SensorReading):
    tensor = torch.tensor([
        reading.temperature,
        reading.vibration,
        reading.pressure,
        reading.power
    ])

    with torch.no_grad():
        error = model.get_reconstruction_error(tensor).item()

    return {
        "is_anomaly": error > threshold,
        "anomaly_score": error,
        "threshold": threshold
    }

```

**Deliverable:** Working API at <http://localhost:8000/docs>

## Month 2 Checkpoint ✓

- PyTorch autoencoder model
- Training pipeline with validation
- Anomaly threshold optimization
- FastAPI inference service
- Docker container for model
- Integration with data pipeline

**Skills Acquired:** - ✓ PyTorch fundamentals (THE critical gap!) - ✓ Custom Dataset/DataLoader - ✓ Training loops - ✓ Model evaluation - ✓ FastAPI ML serving

## MONTH 3: MLOps with MLflow

### Goal: Add experiment tracking and model registry

#### Week 9: MLflow Setup

**Tasks:** - Add MLflow to Docker Compose - Integrate tracking into training script - Log parameters, metrics, and artifacts

**docker-compose addition:**

```

mlflow:
  image: ghcr.io/mlflow/mlflow:v2.9.2
  ports:
    - "5000:5000"
  command: mlflow server --host 0.0.0.0 --port 5000
  volumes:
    - mlflow-data:/mlflow

```

**Training with MLflow:**

```

import mlflow
import mlflow.pytorch

mlflow.set_tracking_uri("http://localhost:5000")

```

```

mlflow.set_experiment("anomaly-detection")

with mlflow.start_run():
    # Log parameters
    mlflow.log_params({
        "input_dim": 4,
        "latent_dim": 2,
        "learning_rate": 1e-3,
        "epochs": 100
    })

    # Training loop
    for epoch in range(epochs):
        train_loss = train_epoch(model, train_loader)
        val_loss = evaluate(model, val_loader)

        mlflow.log_metrics({
            "train_loss": train_loss,
            "val_loss": val_loss
        }, step=epoch)

    # Log model
    mlflow.pytorch.log_model(model, "model")

    # Log threshold
    mlflow.log_metric("optimal_threshold", threshold)
    mlflow.log_metric("f1_score", f1)

```

**Deliverable:** MLflow UI showing experiments at <http://localhost:5000>

## Week 10: Model Registry

**Tasks:** - Register best models in MLflow Registry - Implement model versioning - Create model promotion workflow (staging → production)

**Code:**

```

# mlops/model_registry.py
from mlflow.tracking import MlflowClient

client = MlflowClient()

# Register model
model_uri = f"runs:{run_id}/model"
mv = mlflow.register_model(model_uri, "anomaly-detector")

# Transition to staging
client.transition_model_version_stage(
    name="anomaly-detector",
    version=mv.version,
    stage="Staging"
)

# After validation, promote to production
client.transition_model_version_stage(
    name="anomaly-detector",
    version=mv.version,
    stage="Production"
)

```

## Week 11: DVC for Data Versioning

**Tasks:** - Set up DVC for data versioning - Version training datasets - Create reproducible training pipelines

**Commands:**

```

# Initialize DVC
dvc init

# Track data
dvc add data/training_data.parquet

# Create pipeline
dvc.yaml:
stages:
prepare:
cmd: python scripts/prepare_data.py
deps:
- data/raw/
outs:
- data/processed/

train:
cmd: python ml-models/anomaly_detector/train.py
deps:
- data/processed/
- ml-models/anomaly_detector/model.py
outs:
- models/autoencoder.pt

```

```
metrics:  
  - metrics.json:  
    cache: false
```

## Week 12: Automated Retraining

**Tasks:** - Create script that monitors model performance - Trigger retraining when performance degrades - Automate with simple cron job (CI/CD comes later)

**Deliverable:** Working MLOps pipeline with versioning

### Month 3 Checkpoint ✓

- MLflow tracking integrated
- Model registry with staging/production
- DVC for data versioning
- Reproducible training pipeline
- Automated retraining trigger
- Documentation of MLOps workflow

**Skills Acquired:** - ✓ MLflow (experiment tracking, registry) - ✓ DVC (data versioning) - ✓ MLOps pipelines - ✓ Model lifecycle management

## MONTH 4: Advanced Time Series with PyTorch

### Goal: Build a predictive maintenance model using Transformer architecture

#### Week 13-14: Temporal Fusion Transformer

**Tasks:** - Implement attention-based time series model - Handle multi-horizon forecasting - Predict "time to failure"

**Architecture:**

```
# ml-models/predictive/temporal_transformer.py  
import torch  
import torch.nn as nn  
  
class TemporalAttention(nn.Module):  
    def __init__(self, d_model, n_heads):  
        super().__init__()  
        self.attention = nn.MultiheadAttention(d_model, n_heads)  
        self.norm = nn.LayerNorm(d_model)  
  
    def forward(self, x):  
        attn_out, _ = self.attention(x, x, x)  
        return self.norm(x + attn_out)  
  
class PredictiveMaintenanceModel(nn.Module):  
    def __init__(self, input_dim, d_model=64, n_heads=4, n_layers=3):  
        super().__init__()  
  
        self.input_projection = nn.Linear(input_dim, d_model)  
        self.positional_encoding = PositionalEncoding(d_model)  
  
        self.transformer_layers = nn.ModuleList([  
            TemporalAttention(d_model, n_heads)  
            for _ in range(n_layers)  
        ])  
  
        self.output_layer = nn.Linear(d_model, 1) # Time to failure  
  
    def forward(self, x):  
        # x shape: (batch, seq_len, input_dim)  
        x = self.input_projection(x)  
        x = self.positional_encoding(x)  
  
        for layer in self.transformer_layers:  
            x = layer(x)  
  
        # Take last timestep  
        return self.output_layer(x[:, -1, :])
```

**Resource:** [Temporal Fusion Transformers Paper](#)

#### Week 15-16: Multi-Task Learning

**Tasks:** - Combine anomaly detection + failure prediction - Share encoder, separate heads - Joint training

**Code:**

```
class MultiTaskModel(nn.Module):  
    def __init__(self, input_dim):  
        super().__init__()  
  
        # Shared encoder
```

```

self.encoder = SharedEncoder(input_dim)

# Task-specific heads
self.anomaly_head = AnomalyHead()
self.prediction_head = PredictionHead()

def forward(self, x):
    features = self.encoder(x)

    anomaly_score = self.anomaly_head(features)
    time_to_failure = self.prediction_head(features)

    return {
        "anomaly_score": anomaly_score,
        "time_to_failure": time_to_failure
    }

```

## Month 4 Checkpoint ✓

- Transformer-based time series model
- Positional encoding implementation
- Multi-horizon forecasting
- Multi-task learning architecture
- Model comparison in MLflow

**Skills Acquired:** - ✓ Advanced PyTorch (Transformers!) - ✓ Attention mechanisms - ✓ Time series deep learning - ✓ Multi-task learning

## MONTH 5: Knowledge Graph Integration

### Goal: Build equipment relationship graph with Neo4j

#### Week 17-18: Graph Schema Design

**Tasks:** - Design ontology for industrial equipment - Model: Equipment → Components → Sensors → Readings - Add relationships: DEPENDS\_ON, CONNECTED\_TO, UPSTREAM\_OF

#### Cypher Schema:

```

// Node types
CREATE CONSTRAINT equipment_id IF NOT EXISTS FOR (e:Equipment) REQUIRE e.id IS UNIQUE;
CREATE CONSTRAINT sensor_id IF NOT EXISTS FOR (s:Sensor) REQUIRE s.id IS UNIQUE;

// Equipment node
CREATE (e:Equipment {
    id: "MACHINE_001",
    type: "CNC_Mill",
    manufacturer: "Siemens",
    install_date: date("2020-01-15"),
    location: "Building_A_Floor_2"
})

// Sensor nodes
CREATE (s:Sensor {
    id: "TEMP_001",
    type: "temperature",
    unit: "celsius",
    min_threshold: 20,
    max_threshold: 80
})

// Relationships
MATCH (e:Equipment {id: "MACHINE_001"})
MATCH (s:Sensor {id: "TEMP_001"})
CREATE (e)-[:HAS_SENSOR {position: "spindle"}]->(s)

// Equipment dependencies
MATCH (e1:Equipment {id: "MACHINE_001"})
MATCH (e2:Equipment {id: "MACHINE_002"})
CREATE (e1)-[:FEEDS_INTO {product: "part_A"}]->(e2)

```

#### Week 19-20: Graph-Enhanced ML

**Tasks:** - Use graph features in ML models - Propagate anomalies through connected equipment - Implement graph neural network (optional advanced)

#### Code:

```

# knowledge-graph/graph_features.py
from neo4j import GraphDatabase

class GraphFeatureExtractor:
    def __init__(self, uri, user, password):
        self.driver = GraphDatabase.driver(uri, auth=(user, password))

```

```

def get_equipment_context(self, equipment_id: str) -> dict:
    """Get graph-based features for equipment"""
    query = """
    MATCH (e:Equipment {id: $equipment_id})
    OPTIONAL MATCH (e)-[:DEPENDS_ON]->(upstream:Equipment)
    OPTIONAL MATCH (downstream:Equipment)-[:DEPENDS_ON]->(e)
    OPTIONAL MATCH (e)-[:HAS_SENSOR]->(s:Sensor)
    RETURN
        e.type as equipment_type,
        count(DISTINCT upstream) as upstream_count,
        count(DISTINCT downstream) as downstream_count,
        count(DISTINCT s) as sensor_count,
        collect(DISTINCT upstream.id) as upstream_ids
    """
    with self.driver.session() as session:
        result = session.run(query, equipment_id=equipment_id)
        return result.single().data()

def propagate_anomaly(self, equipment_id: str, anomaly_score: float):
    """Alert downstream equipment of potential issues"""
    query = """
    MATCH (e:Equipment {id: $equipment_id})
    MATCH (downstream:Equipment)-[:DEPENDS_ON*1..3]->(e)
    RETURN downstream.id as affected_equipment,
           length(path) as distance
    """
    # Implementation...

```

## Month 5 Checkpoint ✓

- Neo4J schema for industrial equipment
- Graph population from metadata
- Graph-based feature extraction
- Anomaly propagation algorithm
- Graph visualization dashboard

**Skills Acquired:** - ✓ Advanced Neo4J (beyond certification) - ✓ Graph data modeling - ✓ Graph algorithms for ML - ✓ Cypher optimization

## MONTH 6: RAG System for Technicians

### Goal: Build Q&A system over equipment manuals and logs

#### Week 21-22: Document Processing Pipeline

**Tasks:** - Ingest PDF manuals, maintenance logs - Chunk documents intelligently - Create embeddings with sentence-transformers

**Code:**

```

# llm-assistant/document_processor.py
from langchain.document_loaders import PyPDFLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain.embeddings import HuggingFaceEmbeddings
from langchain.vectorstores import Chroma

class DocumentProcessor:
    def __init__(self):
        self.embeddings = HuggingFaceEmbeddings(
            model_name="sentence-transformers/all-MiniLM-L6-v2"
        )
        self.splitter = RecursiveCharacterTextSplitter(
            chunk_size=500,
            chunk_overlap=50,
            separators=["\n\n", "\n", ".", " "]
        )
        self.vectorstore = Chroma(
            persist_directory="./chroma_db",
            embedding_function=self.embeddings
        )

    def ingest_manual(self, pdf_path: str, equipment_id: str):
        loader = PyPDFLoader(pdf_path)
        documents = loader.load()

        # Add metadata
        for doc in documents:
            doc.metadata["equipment_id"] = equipment_id
            doc.metadata["source_type"] = "manual"

        chunks = self.splitter.split_documents(documents)
        self.vectorstore.add_documents(chunks)

```

#### Week 23-24: RAG Chain with Graph Context

**Tasks:** - Combine vector search with graph context - Create hybrid retrieval - Build conversational chain

**Code:**

```
# llm-assistant/rag_chain.py
from langchain.chains import ConversationalRetrievalChain
from langchain.chat_models import ChatOpenAI
from langchain.prompts import PromptTemplate

class IndustrialRAG:
    def __init__(self, vectorstore, graph_extractor, llm=None):
        self.vectorstore = vectorstore
        self.graph = graph_extractor
        self.llm = llm or ChatOpenAI(model="gpt-3.5-turbo")

        self.prompt = PromptTemplate(
            template="""You are an expert industrial maintenance assistant.

Equipment Context from Knowledge Graph:
{graph_context}

Relevant Documentation:
{documents}

Current Sensor Readings:
{sensor_data}

User Question: {question}

Provide a helpful, technically accurate response:""",
            input_variables=["graph_context", "documents", "sensor_data", "question"]
        )

    def query(self, question: str, equipment_id: str) -> str:
        # Get graph context
        graph_context = self.graph.get_equipment_context(equipment_id)

        # Get relevant documents
        docs = self.vectorstore.similarity_search(
            question,
            k=3,
            filter={"equipment_id": equipment_id}
        )

        # Get current sensor data
        sensor_data = self.get_current_readings(equipment_id)

        # Generate response
        response = self.llm.invoke(
            self.prompt.format(
                graph_context=graph_context,
                documents=docs,
                sensor_data=sensor_data,
                question=question
            )
        )

        return response.content
```

## Month 6 Checkpoint ✓

- Document ingestion pipeline
- ChromaDB vector store
- Hybrid retrieval (vector + graph)
- Conversational RAG chain
- Chat UI in Streamlit
- Evaluation on test questions

**Skills Acquired:** - ✓ Advanced RAG architectures - ✓ Hybrid retrieval systems - ✓ Vector databases (ChromaDB) - ✓ LangChain advanced patterns

## MONTH 7: LLM Fine-tuning

### Goal: Fine-tune open-source LLM for industrial domain

#### Week 25-26: Dataset Preparation

**Tasks:** - Create instruction-tuning dataset from maintenance logs - Format: (instruction, input, output) triplets - Quality filtering and deduplication

**Dataset Format:**

```
{
  "instruction": "Diagnose the issue based on the sensor readings",
  "input": "Equipment: CNC_Mill_001\nTemperature: 85°C (normal: 40-70°C)\nVibration: 2.5mm/s (normal: 0.5-1.5mm/s)\nPressure: Normal",
  "output": "The high temperature and elevated vibration suggest bearing wear. Recommended actions:\n1. Inspect spindle bearings\n2. Ch
```

## Week 27-28: LoRA Fine-tuning

**Tasks:** - Fine-tune Mistral-7B or LLaMA-2-7B with LoRA - Use QLoRA for memory efficiency - Implement evaluation metrics

**Code:**

```
# llm-assistant/finetune.py
from transformers import AutoModelForCausalLM, AutoTokenizer, BitsAndBytesConfig
from peft import LoraConfig, get_peft_model, prepare_model_for_kbit_training
from datasets import load_dataset
from trl import SFTTrainer

# Quantization config for QLoRA
bnb_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_quant_type="nf4",
    bnb_4bit_compute_dtype=torch.float16,
    bnb_4bit_use_double_quant=True
)

# Load model
model = AutoModelForCausalLM.from_pretrained(
    "mistralai/Mistral-7B-v0.1",
    quantization_config=bnb_config,
    device_map="auto"
)

# LoRA config
lora_config = LoraConfig(
    r=16, # rank
    lora_alpha=32,
    target_modules=["q_proj", "k_proj", "v_proj", "o_proj"],
    lora_dropout=0.05,
    bias="none",
    task_type="CAUSAL_LM"
)

# Prepare model
model = prepare_model_for_kbit_training(model)
model = get_peft_model(model, lora_config)

# Train
trainer = SFTTrainer(
    model=model,
    train_dataset=dataset,
    dataset_text_field="text",
    max_seq_length=512,
    args=TrainingArguments(
        output_dir=".//industrial-mistral-lora",
        per_device_train_batch_size=4,
        gradient_accumulation_steps=4,
        num_train_epochs=3,
        learning_rate=2e-4,
        fp16=True,
        logging_steps=10,
        save_strategy="epoch"
    )
)
trainer.train()
```

**Resource:** [Hugging Face PEFT Guide \(FREE\)](#)

## Month 7 Checkpoint ✓

- Instruction-tuning dataset created
- LoRA fine-tuning pipeline
- Model evaluation (perplexity, task accuracy)
- Merged LoRA weights for inference
- Comparison: base vs fine-tuned
- Model uploaded to Hugging Face Hub

**Skills Acquired:** - ✓ LLM Fine-tuning (LoRA, QLoRA, PEFT) - ✓ Instruction tuning - ✓ Quantization techniques - ✓ Hugging Face ecosystem

## MONTH 8: Kubernetes Deployment

### Goal: Deploy entire system on Kubernetes

#### Week 29-30: Kubernetes Fundamentals

**Tasks:** - Set up local K8s (minikube or kind) - Create deployments for each service - Configure services and ingress

**Learn While Building:**

```

# kubernetes/deployments/ml-api.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: ml-api
  labels:
    app: ml-api
spec:
  replicas: 2
  selector:
    matchLabels:
      app: ml-api
  template:
    metadata:
      labels:
        app: ml-api
    spec:
      containers:
        - name: ml-api
          image: industrialmind/ml-api:latest
          ports:
            - containerPort: 8000
          resources:
            requests:
              memory: "512Mi"
              cpu: "250m"
            limits:
              memory: "1Gi"
              cpu: "500m"
      env:
        - name: MLFLOW_TRACKING_URI
          value: "http://mlflow:5000"
        - name: MODEL_NAME
          value: "anomaly-detector"
        - name: MODEL_STAGE
          value: "Production"
---
apiVersion: v1
kind: Service
metadata:
  name: ml-api
spec:
  selector:
    app: ml-api
  ports:
    - port: 80
      targetPort: 8000
  type: ClusterIP

```

## Week 31-32: Helm Charts & Scaling

**Tasks:** - Create Helm chart for easy deployment - Implement horizontal pod autoscaling - Set up GPU node pool for inference

**Helm Chart Structure:**

```

industrialmind-chart/
├── Chart.yaml
├── values.yaml
└── templates/
    ├── deployment.yaml
    ├── service.yaml
    ├── ingress.yaml
    └── configmap.yaml
        └── hpa.yaml

```

**Resource:** [Kubernetes in 1 Hour](#) (FREE)

## Month 8 Checkpoint ✓

- All services running on K8s
- Helm chart for deployment
- Horizontal Pod Autoscaler
- Ingress configuration
- Secrets management
- Local cluster fully functional

**Skills Acquired:** - ✓ Kubernetes fundamentals - ✓ Helm charts - ✓ Container orchestration - ✓ Scaling strategies

## MONTH 9: Cloud Deployment (AWS)

**Goal:** Deploy to AWS with SageMaker integration

**Week 33-34: AWS Infrastructure**

**Tasks:** - Set up EKS cluster - Deploy to AWS with Terraform - Integrate with AWS services

#### Terraform:

```
# terraform/main.tf
provider "aws" {
  region = "eu-west-1"
}

module "eks" {
  source      = "terraform-aws-modules/eks/aws"
  cluster_name = "industrialmind-cluster"
  cluster_version = "1.28"

  vpc_id     = module.vpc.vpc_id
  subnet_ids = module.vpc.private_subnets

  eks_managed_node_groups = {
    general = {
      desired_size = 2
      min_size     = 1
      max_size     = 4
      instance_types = ["t3.medium"]
    }

    ml = {
      desired_size = 1
      min_size     = 0
      max_size     = 2
      instance_types = ["g4dn.xlarge"] # GPU
      labels = {
        workload = "ml-inference"
      }
    }
  }
}
```

## Week 35-36: SageMaker Integration

**Tasks:** - Deploy models to SageMaker endpoints - Set up SageMaker Pipelines - Implement A/B testing

#### Code:

```
# mlops/sagemaker_deploy.py
import sagemaker
from sagemaker.pytorch import PyTorchModel

def deploy_to_sagemaker(model_artifact_path: str):
    role = sagemaker.get_execution_role()

    pytorch_model = PyTorchModel(
        model_data=model_artifact_path,
        role=role,
        framework_version="2.0",
        py_version="py310",
        entry_point="inference.py"
    )

    predictor = pytorch_model.deploy(
        instance_type="ml.g4dn.xlarge",
        initial_instance_count=1,
        endpoint_name="anomaly-detector-prod"
    )

    return predictor
```

## Month 9 Checkpoint ✓

- EKS cluster running
- Terraform infrastructure as code
- SageMaker endpoints deployed
- Cost monitoring set up
- AWS Well-Architected review
- Pass AWS ML Specialty exam!

**Skills Acquired:** - ✓ AWS EKS - ✓ AWS SageMaker - ✓ Terraform - ✓ Cloud architecture

## MONTH 10: CI/CD Pipeline

### Goal: Fully automated deployment pipeline

#### Week 37-38: GitHub Actions for ML

**Tasks:** - Automated testing on PR - Model training on merge to main - Automated deployment to staging

### .github/workflows/ml-pipeline.yaml:

```
name: ML Pipeline

on:
  push:
    branches: [main]
    paths:
      - 'ml-models/**'
      - 'data/**'
  pull_request:
    branches: [main]

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4

      - name: Set up Python
        uses: actions/setup-python@v4
        with:
          python-version: '3.10'

      - name: Install dependencies
        run: pip install -r requirements.txt

      - name: Run tests
        run: pytest tests/ -v

      - name: Run model tests
        run: python -m pytest ml-models/tests/ -v

  train:
    needs: test
    if: github.ref == 'refs/heads/main'
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4

      - name: Configure AWS credentials
        uses: aws-actions/configure-aws-credentials@v4
        with:
          aws-access-key-id: ${{ secrets.AWS_ACCESS_KEY_ID }}
          aws-secret-access-key: ${{ secrets.AWS_SECRET_ACCESS_KEY }}
          aws-region: eu-west-1

      - name: Train model
        run: |
          python ml-models/anomaly_detector/train.py \
            --mlflow-tracking-uri ${{ secrets.MLFLOW_URI }}

      - name: Register model
        run: python mlops/register_model.py

  deploy-staging:
    needs: train
    runs-on: ubuntu-latest
    environment: staging
    steps:
      - name: Deploy to staging
        run: |
          aws sagemaker update-endpoint \
            --endpoint-name anomaly-detector-staging \
            --endpoint-config-name ${{ steps.train.outputs.config_name }}
```

## Week 39-40: Model Validation Gates

**Tasks:** - Automated model validation before production - Performance regression checks - Data drift detection

**Code:**

```
# mlops/validation_gate.py
class ModelValidationGate:
    def __init__(self, production_model, candidate_model):
        self.prod = production_model
        self.candidate = candidate_model

    def validate(self, test_data) -> bool:
        """Check if candidate model passes all gates"""

        # Performance gate
        prod_metrics = self.evaluate(self.prod, test_data)
        candidate_metrics = self.evaluate(self.candidate, test_data)

        if candidate_metrics['f1'] < prod_metrics['f1'] * 0.95:
            return False, "F1 score regression > 5%"
```

```

# Latency gate
if candidate_metrics['p99_latency'] > 100: # ms
    return False, "P99 latency exceeds 100ms"

# Bias gate
if self.detect_bias(self.candidate, test_data):
    return False, "Model shows significant bias"

return True, "All gates passed"

```

## Month 10 Checkpoint ✓

- GitHub Actions CI/CD pipeline
- Automated testing
- Automated training
- Model validation gates
- Staging deployment automation
- Production deployment with approval

**Skills Acquired:** - ✓ CI/CD for ML - ✓ GitHub Actions - ✓ Automated testing - ✓ Deployment automation

## MONTH 11: Monitoring & Observability

### Goal: Production-grade monitoring

#### Week 41-42: Prometheus & Grafana

**Tasks:** - Model performance monitoring - Data drift detection - Alert configuration

##### Metrics to Track:

```

# api/metrics.py
from prometheus_client import Counter, Histogram, Gauge

# Request metrics
PREDICTION_COUNTER = Counter(
    'predictions_total',
    'Total predictions',
    ['model_version', 'result']
)

PREDICTION_LATENCY = Histogram(
    'prediction_latency_seconds',
    'Prediction latency',
    buckets=[0.01, 0.025, 0.05, 0.1, 0.25, 0.5, 1.0]
)

# Model metrics
ANOMALY_SCORE = Histogram(
    'anomaly_score',
    'Distribution of anomaly scores',
    buckets=[0.01, 0.05, 0.1, 0.2, 0.5, 1.0]
)

DATA_DRIFT_SCORE = Gauge(
    'data_drift_score',
    'Current data drift score',
    ['feature']
)

```

#### Week 43-44: Data Drift & Model Decay

**Tasks:** - Implement drift detection - Set up alerting - Create Grafana dashboards

##### Code:

```

# mlops/drift_detection.py
from evidently.metrics import DataDriftTable
from evidently.report import Report

class DriftDetector:
    def __init__(self, reference_data):
        self.reference = reference_data

    def check_drift(self, current_data) -> dict:
        report = Report(metrics=[DataDriftTable()])
        report.run(
            reference_data=self.reference,
            current_data=current_data
        )

        results = report.as_dict()

        # Update Prometheus metrics

```

```

        for feature, drift_score in results['drift_by_feature'].items():
            DATA_DRIFT_SCORE.labels(feature=feature).set(drift_score)

    return results

```

## Month 11 Checkpoint ✓

- Prometheus metrics collection
- Grafana dashboards
- Data drift detection
- Model performance monitoring
- Alerting configuration
- Runbook documentation

**Skills Acquired:** - ✓ Prometheus & Grafana - ✓ ML monitoring - ✓ Data drift detection - ✓ Observability best practices

## MONTH 12: Polish & Job Search

### Goal: Portfolio-ready project + Active job search

#### Week 45-46: Documentation & Demo

**Tasks:** - Complete README with architecture diagram - Create demo video (5-10 minutes) - Write blog post about the project - Prepare technical presentation

##### README Structure:

```

# IndustrialMind 🚀

> End-to-end ML platform for industrial predictive maintenance

[![CI/CD](badge)](link)
[![License](badge)](link)
[![Demo](badge)](link)

## 📺 What This Project Demonstrates

- **PyTorch** deep learning for anomaly detection
- **Transformer** architecture for time series
- **MLOps** pipeline with MLflow, DVC
- **Knowledge Graph** with Neo4J
- **LLM/RAG** system with fine-tuned model
- **Kubernetes** deployment at scale
- **AWS SageMaker** integration

## 🏗️ Architecture

[Architecture diagram]

## 🚀 Quick Start

```bash
docker-compose up -d
```

```

## Results

| Model       | F1 Score | Latency (p99) |
|-------------|----------|---------------|
| Autoencoder | 0.92     | 15ms          |
| Transformer | 0.95     | 45ms          |

[More documentation...]

```

### Week 47-48: Active Job Search

**Tasks:**
- Apply to 50+ positions
- Tailor applications to each role
- Prepare for technical interviews

**Application Tracker:**

Company	Role	Location	Status	Notes
Google	ML Engineer	Zurich	Applied	
Roche	AI Lead	Basel	Applied	
G42	Senior ML	Dubai	Applied	
...	...	...	...	

### Month 12 Checkpoint ✓

```

README polished  Demo video recorded  Blog post published  50+ applications sent  Interview prep complete  First interviews scheduled!

```
---
```

```
## 📊 SKILLS MATRIX: Before vs After
```

| Skill           | Before | After | Evidence            |
|-----------------|--------|-------|---------------------|
| PyTorch         | **     | ***** | 3 production models |
| MLOps           | **     | ***** | Full pipeline       |
| Kubernetes      | **     | ****  | EKS deployment      |
| AWS             | *      | ****  | SageMaker + cert    |
| LLM Fine-tuning | **     | ***** | Published model     |
| CI/CD           | ***    | ***** | GitHub Actions      |
| Monitoring      | **     | ****  | Prometheus/Grafana  |

```
---
```

```
## 💡 EXPECTED OUTCOMES
```

### ### Portfolio Impact

After completing IndustrialMind, you'll have:

- \*\*GitHub repo\*\* with 1000+ commits, production-quality code
- \*\*Published model\*\* on Hugging Face Hub
- \*\*Blog posts\*\* showing thought leadership
- \*\*Demo video\*\* for quick showcasing
- \*\*Live deployment\*\* (even if small scale)

### ### Interview Talking Points

Every component maps to interview questions:

| Question                          | Your Answer                            |
|-----------------------------------|--|
| "Tell me about a PyTorch project" | IndustrialMind anomaly detector        |
| "How do you handle MLOps?"        | MLflow + DVC + GitHub Actions          |
| "Experience with Kubernetes?"     | Deployed entire platform on EKS        |
| "LLM experience?"                 | Fine-tuned Mistral with LoRA           |
| "How do you monitor models?"      | Prometheus + Grafana + drift detection |

### ### Salary Expectation

With this portfolio, you can confidently target:

| Region      | Role               | Expected Salary       |
|-------------|--------------------|-----------------------|
| Switzerland | Senior ML Engineer | CHF 140-170K          |
| UAE         | ML Lead            | AED 500-650K (0% tax) |
| UK          | Staff ML Engineer  | £100-130K             |
| USA         | Senior MLE         | \$180-250K            |

```
---
```

```
## 🚀 START NOW
```

### ### This Week's Tasks

Create GitHub repo: industrialmind  Set up basic project structure  Write initial README  Install Docker and Docker Compose  Start Week 1 tasks

## Resources You Need

| Resource        | Purpose           | Cost |
|-----------------|-------------------|------|
| GitHub          | Code hosting      | FREE |
| Docker Desktop  | Containers        | FREE |
| AWS Free Tier   | Cloud (12 months) | FREE |
| Hugging Face    | Models            | FREE |
| MLflow          | Tracking          | FREE |
| Neo4j Community | Graph DB          | FREE |

**Total Cost: €0** (using free tiers)

## ✍ ACCOUNTABILITY

## Weekly Check-in Template

Every Sunday, answer:

1. ✓ What did I ship this week?
2. ☀ What's the goal for next week?
3. 🚫 What's blocking me?
4. 💡 What did I learn?

## Monthly Milestones

| Month | Must Ship          | Nice to Have             |
|-------|--------------------|--------------------------|
| 1     | Data pipeline      | Performance optimization |
| 2     | PyTorch model      | Advanced architecture    |
| 3     | MLflow integration | Automated retraining     |
| 4     | Transformer model  | Multi-task learning      |
| 5     | Knowledge Graph    | Graph neural network     |
| 6     | RAG system         | Hybrid retrieval         |
| 7     | Fine-tuned LLM     | Multi-modal              |
| 8     | K8s deployment     | Auto-scaling             |
| 9     | AWS deployment     | Multi-region             |
| 10    | CI/CD pipeline     | Canary deployments       |
| 11    | Monitoring         | Automated remediation    |
| 12    | Job offers!        | Multiple offers          |

"The best time to plant a tree was 20 years ago. The second best time is now."

**Let's build.** 🚀